




deeplearning.ai

Comparing to human-
level performance

Why human-level
performance?

Comparing to human-level performance



Why compare to human-level performance

Humans are quite good at a lot of tasks. So long as ML is worse than humans, you can:

- - Get labeled data from humans. (x, y)
- - Gain insight from manual error analysis:
Why did a person get this right?
- - Better analysis of bias/variance.




deeplearning.ai

Comparing to human-
level performance


Avoidable bias

Bias and Variance




high bias

underfitting



“just right”



high variance

overfitting


Bias and Variance

Cat classification

Human-level $\approx 0\%$ ----

Training set error:

Dev set error:



high variance




high bias

high bias
high variance

low bias
low variance

Cat classification example





deeplearning.ai

Comparing to human-level performance

Understanding
human-level
performance

Human-level error as a proxy for Bayes error

Medical image classification example:

Suppose:

(a) Typical human 3 % error

→ (b) Typical doctor 1 % error

(c) Experienced doctor 0.7 % error

→ (d) Team of experienced doctors .. 0.5 % error ←

Bayes error \leq 0.5 %

What is “human-level” error?



Error analysis example


Human (proxy for Bayes error)

↕ Avoidable bias


Training error

↕ Variance


Dev error




↑ Bias



↑ Variance



Summary of bias/variance with human-level performance





deeplearning.ai

Comparing to human-
level performance

Surpassing human-
level performance

Surpassing human-level performance

Team of humans

0.5%

One human

0.1 \updownarrow ~~1.0%~~


Training error

0.6%

Dev error

0.2 \updownarrow 0.8%

What is avoidable bias?



Problems where ML significantly surpasses human-level performance

- - Online advertising
- - Product recommendations
- - Logistics (predicting transit time)
- - Loan approvals

Structured data

Not natural perception

Lots of data

- Speech recognition
- Some image recognition
- Medical
 - ECG, Skin cancer, ...



deeplearning.ai

Comparing to human-
level performance

Improving your model
performance

The two fundamental assumptions of supervised learning

1. You can fit the training set pretty well.



~ Avoidable bias

2. The training set performance generalizes pretty well to the dev/test set.



~ Variance

Reducing (avoidable) bias and variance

Human-level

↑ Avoidable bias
↓

Training error

↑ Variance
↓

Dev error

Train bigger model

Train longer/better optimization algorithms
- momentum, RMSprop, Adam

NN architecture/hyperparameters search RNN
CNN

More data

Regularization

- L_2 , dropout, data augmentation

NN architecture/hyperparameters search